
An abstract graphic on the left side of the slide, consisting of a complex network of blue lines and dots. The dots are of varying sizes and are connected by thin blue lines, creating a web-like structure that tapers towards the right. The overall effect is a sense of depth and connectivity, typical of network diagrams or data visualizations.

Deep Sequence Modeling

Given an image of a ball,
can you predict where it will go next?



Given an image of a ball,
can you predict where it will go next?



Sequences in the Wild



Audio

Sequences in the Wild

Introduction to Deep Learning

Text

A Sequence Modeling Problem: Predict the Next Word

A Sequence Modeling Problem: Predict the Next Word

“This morning I took my cat for a walk.”

given these words

predict the
next word

Idea #1: Use a Fixed Window

“This morning I took my cat for a walk.”

given these predict the
two words next word

One-hot feature encoding: tells us what each word is

[1 0 0 0 0 0 1 0 0 0]

for

a



prediction

Problem #1: Can't Model Long-Term Dependencies

“**France** is where I grew up, but I now live in Boston. I speak fluent ____.”

We need information from **the distant past** to accurately predict the correct word.

Idea #2: Use Entire Sequence as Set of Counts

“This morning I took my cat for a”



“bag of words”

[0 1 0 0 1 0 0 ... 0 0 1 1 0 0 0 1]



prediction

Problem #2: Counts Don't Preserve Order



The food was good, not bad at all.

vs.

The food was bad, not good at all.



Idea #3: Use a Really Big Fixed Window

"This morning I took my cat for a walk."

given these
words

predict the
next word

[1 0 0 0 0 0 0 0 0 0 1 0 0 1 0 0 0 1 0 0 0 0 0 0 0 1 0 ...]

morning

I

took

this

cat



prediction

Problem #3: No Parameter Sharing

[1 0 0 0 0 0 0 0 0 0 1 0 0 1 0 0 0 0 1 0 0 0 0 0 0 0 1 0 ...]

this morning took the cat

Each of these inputs has a **separate** parameter:

[0 0 0 1 0 0 0 1 0 0 0 1 0 0 0 0 1 0 0 0 0 0 0 0 0 0 1 ...]


this morning

Things we learn about the sequence **won't transfer** if they appear **elsewhere** in the sequence.

Sequence Modeling: Design Criteria

To model sequences, we need to:


1. Handle **variable-length** sequences
2. Track **long-term** dependencies
3. Maintain information about **order**
4. **Share parameters** across the sequence




Recurrent Neural Networks (RNNs) as
an approach to sequence modeling problems

Recurrent Neural Networks (RNNs)


Recurrent Neural Networks for Sequence Modeling



One to One
"Vanilla" neural network




Many to One
Sentiment Classification



Many to Many
Music Generation

... and many other
architectures and
applications

Recurrent Neural Network (RNN)



Apply a **recurrence relation** at every time step to process a sequence:

$$\boxed{h_t} = \boxed{f_W}(\boxed{h_{t-1}}, \boxed{x_t})$$

cell state function parameterized by W old state input vector at time step t

Note: the same function and set of parameters are used at every time step


RNN Intuition

```
my_rnn = RNN()
hidden_state = [0, 0, 0, 0]


sentence = ["I", "love", "recurrent", "neural"]

for word in sentence:
    prediction, hidden_state = my_rnn(word, hidden_state)

next_word_prediction = prediction
# >>> "networks!"
```



RNN State Update and Output



Output Vector

$$\hat{y}_t = \mathbf{W}_{hy}^T h_t$$

Update Hidden State



$$h_t = \tanh(\mathbf{W}_{hh}^T h_{t-1} + \mathbf{W}_{xh}^T x_t)$$

Input Vector

$$x_t$$



$$\tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$

Mathematical formula of the Tanh function




ReLU

$$R(z) = \max(0, z)$$




RNNs: Computational Graph Across Time

Re-use the **same weight matrices** at every time step



RNNs: Computational Graph Across Time



RNNs from Scratch



```
class MyRNNCell(tf.keras.layers.Layer):
    def __init__(self, rnn_units, input_dim, output_dim):
        super(MyRNNCell, self).__init__()


        # Initialize weight matrices
        self.W_xh = self.add_weight([rnn_units, input_dim])
        self.W_hh = self.add_weight([rnn_units, rnn_units])
        self.W_hy = self.add_weight([output_dim, rnn_units])

        # Initialize hidden state to zeros
        self.h = tf.zeros([rnn_units, 1])

    def call(self, x):
        # Update the hidden state
        self.h = tf.math.tanh( self.W_hh * self.h + self.W_xh * x )

        # Compute the output
        output = self.W_hy * self.h


        # Return the current output and hidden state
        return output, self.h
```



RNN Implementation in TensorFlow




```
tf.keras.layers.SimpleRNN(rnn_units)
```



Backpropagation Through Time (BPTT)


Recall: Backpropagation in Feed Forward Models




Backpropagation algorithm:

1. Take the derivative (gradient) of the loss with respect to each parameter
2. Shift parameters in order to minimize loss


RNNs: Backpropagation Through Time



RNNs: Backpropagation Through Time



Standard RNN Gradient Flow: Exploding Gradients



Computing the gradient wrt h_0 involves many factors of W_{hh} + repeated gradient computation!

Many values > 1 :
exploding gradients

Gradient clipping to
scale big gradients

Many values < 1 :
vanishing gradients

1. Activation function
2. Weight initialization
3. Network architecture

The Problem of Long-Term Dependencies

Why are vanishing gradients a problem?

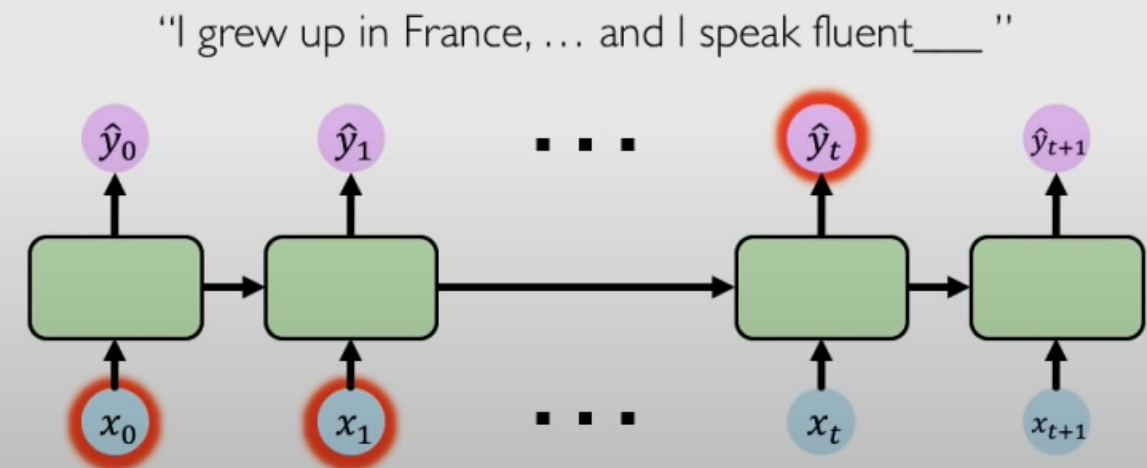

Multiply many **small numbers** together




Errors due to further back time steps
have smaller and smaller gradients




Bias parameters to capture short-term
dependencies




Without gradient clipping



With gradient clipping



Trick #1: Activation Functions



Using ReLU prevents f' from shrinking the gradients when $x > 0$

Trick #2: Parameter Initialization

Initialize **weights** to identity matrix

Initialize **biases** to zero

$$I_n = \begin{pmatrix} 1 & 0 & 0 & \dots & 0 \\ 0 & 1 & 0 & \dots & 0 \\ 0 & 0 & 1 & \dots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \dots & 1 \end{pmatrix}$$

This helps prevent the weights from shrinking to zero.